



DECISION TREE

Anurag Nagar

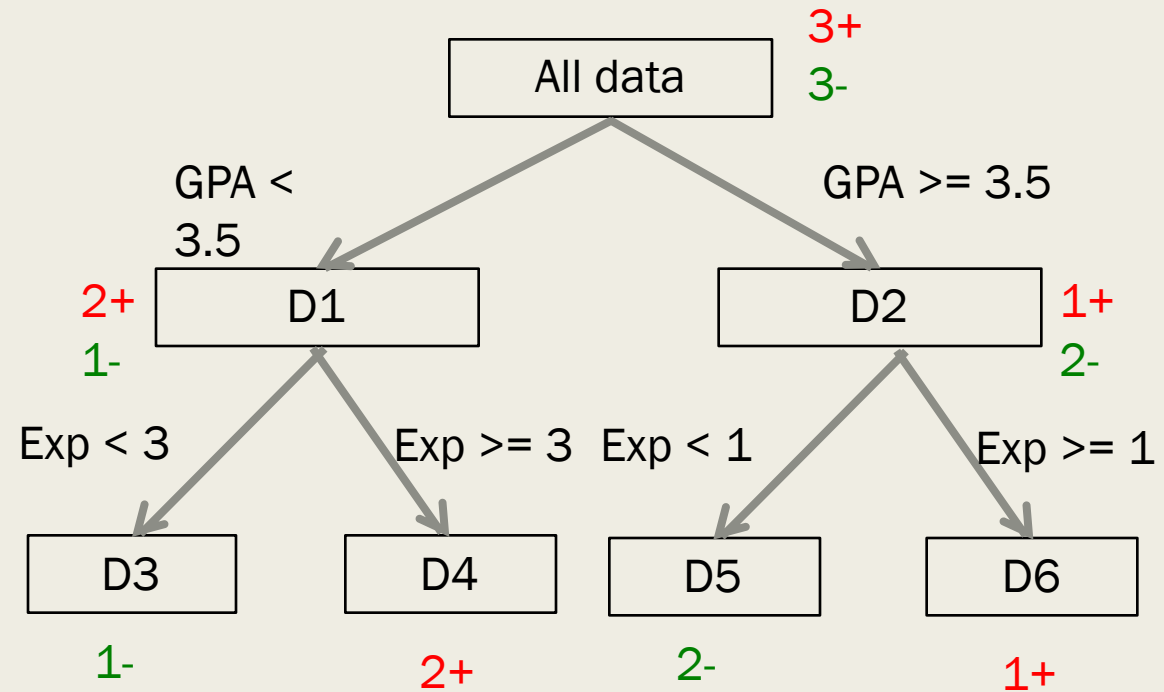


What is a DT?

Decision Tree

■ 2-D case

GPA	Years Exp	Internship
4.0	0	0
3.0	1	0
3.4	4	1
3.6	2	0
3.8	4	1
2.5	3	1



Of course, you could start with Exp as the first sorting or splitting criteria and get a different tree.

Decision Tree - Representation

- In this class, DT is a way to represent concepts & hypothesis about a target concept
- Can be written in form of **rules**

IF exp > 3 and GPA > 3.5 THEN internship = 1


- **Leaf nodes** decide values of output variable
- **Internal nodes & edges** represent splitting (sorting) criteria.
- We will consider Boolean attributes and output.
- Given instances with n Boolean attributes i.e. each X^i is of the form:

$X^i = (x_1, x_2, \dots, x_n)$ e.g. $X^i = (0, 1, \dots, 0)$

How will you represent one hypothesis

=> A binary tree

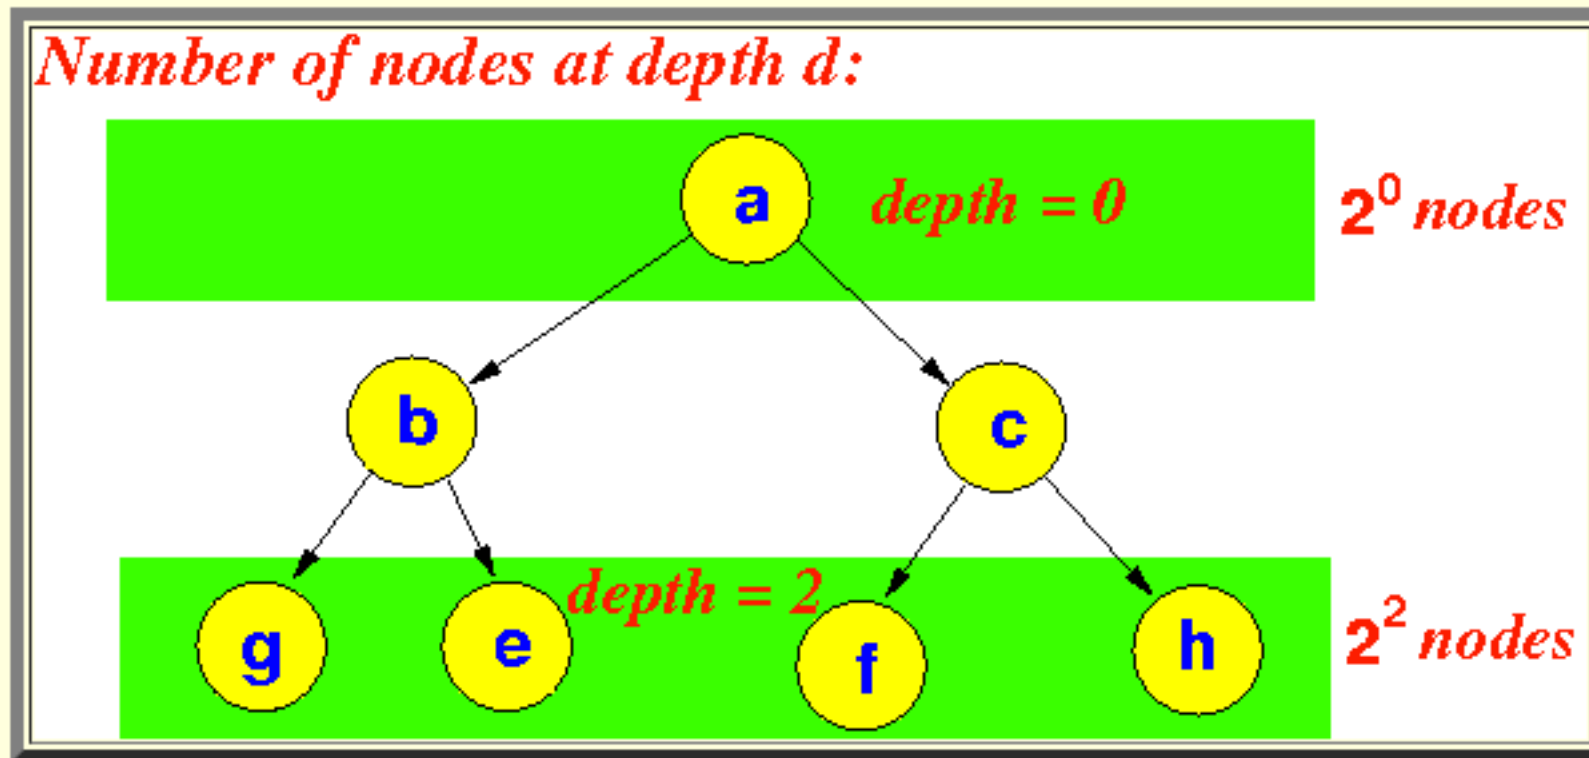
X1	X2	X3
0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0
1	1	1



Decision tree can be used to represent instances with edges representing sorting conditions


Properties of Binary Trees

- A complete binary tree of height h has $2^{h+1} - 1$ nodes
- Number of nodes at depth d is 2^d



Properties of Binary Trees

- A complete binary tree of height h has $2^{h+1} - 1$ nodes
- Number of nodes at depth d is 2^d



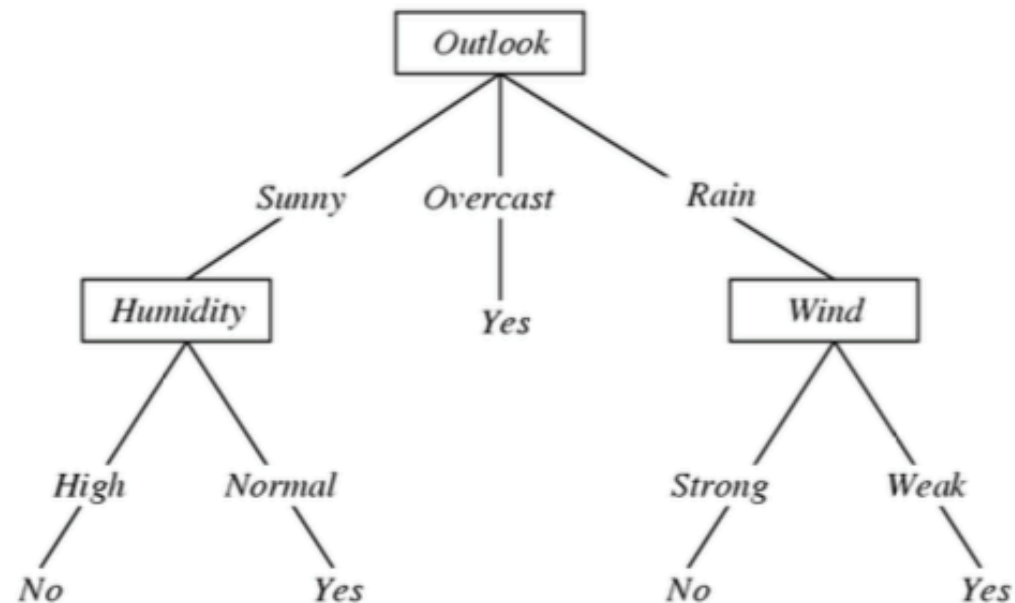
$$\# \text{ nodes} = 2^0 + 2^1 + \dots + 2^h = 2^{h+1} - 1$$

Learning a DT

Use of DT in learning

- Use the training examples and their labels to construct decision tree
- For example, (X^1, y^1) could be $((0, 0, 0), 1)$
- You can use DT to model knowledge from training data.

A Decision tree for
 $F: \langle \text{Outlook, Humidity, Wind, Temp} \rangle \rightarrow \text{PlayTennis?}$



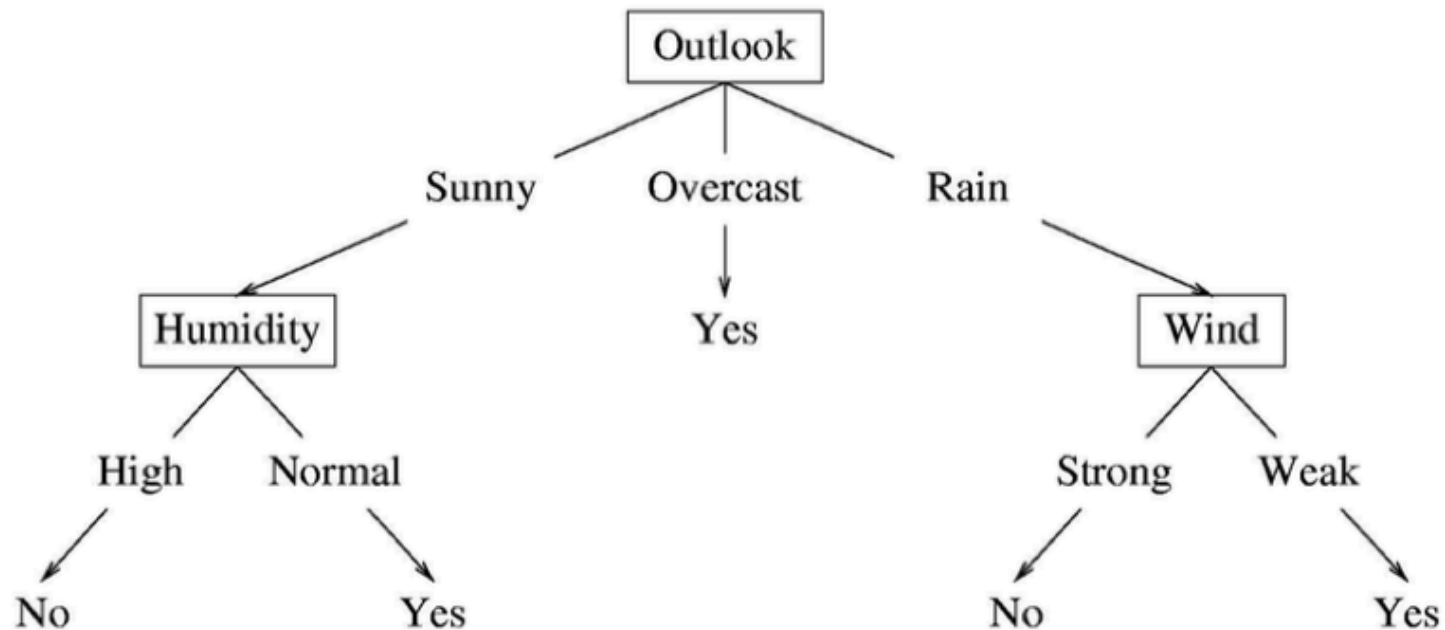
Each internal node: test one discrete-valued attribute X_i

Each branch from a node: selects one value for X_i

Each leaf node: predict Y (or $P(Y|X \in \text{leaf})$)

Using DT to represent hypotheses



- **Internal nodes** test the value of particular features x_j and branch according to the results of the test.
- **Leaf nodes** specify the class $h(\mathbf{x})$.



Suppose the features are **Outlook** (x_1), **Temperature** (x_2), **Humidity** (x_3), and **Wind** (x_4). Then the feature vector $\mathbf{x} = (\text{Sunny}, \text{Hot}, \text{High}, \text{Strong})$ will be classified as **No**. The **Temperature** feature is irrelevant.



Classification Boundary of a DT

Decision trees divide the feature space into axis-parallel rectangles, and label each rectangle with one of the K classes.





Classification Boundary of a DT

Decision trees divide the feature space into axis-parallel rectangles, and label each rectangle with one of the K classes.





Classification Boundary of a DT

Decision trees divide the feature space into axis-parallel rectangles, and label each rectangle with one of the K classes.





Classification Boundary of a DT

Decision trees divide the feature space into axis-parallel rectangles, and label each rectangle with one of the K classes.



Classification Boundary of a DT

Decision trees divide the feature space into axis-parallel rectangles, and label each rectangle with one of the K classes.



Hypothesis Space of DT

As the number of nodes (or depth) of tree increases, the hypothesis space grows

- **depth 1** (“decision stump”) can represent any boolean function of one feature.
- **depth 2** Any boolean function of two features; some boolean functions involving three features (e.g., $(x_1 \wedge x_2) \vee (\neg x_1 \wedge \neg x_3)$)
- **etc.**


Finding the best split (also called sort)

DT – How to find best sorting?


- Which attribute should I sort (split) on first?
 - It DOES make a difference.
- Informally, we want that split that gives maximum purity at each node i.e. split such that all instances are of a single class (or close to it).

Which test is more informative?

**Split over whether
Balance exceeds 50K**



**Split over whether
applicant is employed**



Entropy

Entropy is a measure of Information Content (IC).

$$H(X) = \sum -p_i \log_2 p_i$$

where p_i is the probability of the i^{th} class.

If you think deeply, it is the **expected value** of $-\log_2 p_i$ or $\log(1/p_i)$.


This quantity is also known as information of an attribute.

Also, $H(x)$ can be thought of as the number of bits needed to encode a dataset

- Entropy = $\sum_i -p_i \log_2 p_i$

p_i is the probability of class i

Compute it as the proportion of class i in the set.



16/30 are green circles; 14/30 are pink crosses

$\log_2(16/30) = -.9$; $\log_2(14/30) = -1.1$

Entropy = $-(16/30)(-.9) - (14/30)(-1.1) = .99$

- Entropy comes from information theory. The higher the entropy the more the information content.


What does that mean for learning from examples?

- What is the entropy of a group in which all examples belong to the same class?

- $\text{entropy} = -1 \log_2 1 = 0$

not a good training set for learning

**Minimum
impurity**




- What is the entropy of a group with 50% in either class?


- $\text{entropy} = -0.5 \log_2 0.5 - 0.5 \log_2 0.5 = 1$

good training set for learning

**Maximum
impurity**



Entropy of a binary random variable



- Entropy is maximum at $p=0.5$
- Entropy is zero and $p=0$ or $p=1$.

Information Gain

- We want to determine **which attribute** in a given set of training feature vectors is **most useful** for discriminating between the classes to be learned.
- **Information gain** tells us how important a given attribute of the feature vectors is.
- We will use it to decide the ordering of attributes in the nodes of a decision tree.

Information Gain

- Suppose you just know the class labels initially.
- Then you know one of the attributes.

=> Does it really help you?

=> Do you get any information gain or reduction in entropy ?

=> Do you get any increase in purity of the classes by knowing an attribute?



Mean the same thing

Example of IG

Predicting credit risk

<2 years at current job?	missed payments?	defaulted?
N	N	N
Y	N	Y
N	N	N
N	N	N
N	Y	Y
Y	N	N
N	Y	N
N	Y	Y
Y	N	N
Y	N	N

Class attribute is defaulted?

Independent Attributes - the first two

How many bits does it take to specify the attribute of 'defaulted'?

- $P(\text{defaulted} = Y) = 3/10$
- $P(\text{defaulted} = N) = 7/10$

$$\begin{aligned} H(Y) &= - \sum_{i=Y,N} P(Y = y_i) \log_2 P(Y = y_i) \\ &= -0.3 \log_2 0.3 - 0.7 \log_2 0.7 \\ &= 0.8813 \end{aligned}$$

Can you do better than this by knowing another attribute?

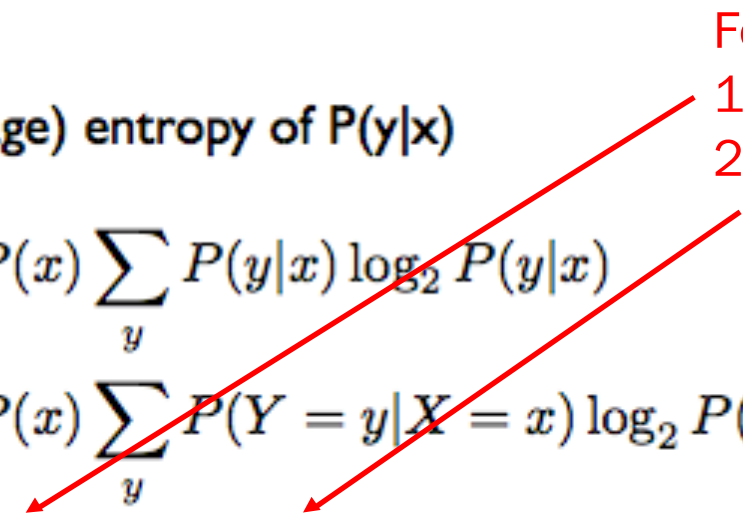
Mutual Information or Information Gain

- Conditional Entropy:
 - How much reduction in entropy (or increase in purity) of class attribute do you get by knowing an **additional variable**

- $H(Y|X)$ is the remaining entropy of Y given X

or

The expected (or average) entropy of $P(y|x)$

$$\begin{aligned} H(Y|X) &\equiv - \sum_x P(x) \sum_y P(y|x) \log_2 P(y|x) \\ &= - \sum_x P(x) \sum_y P(Y = y|X = x) \log_2 P(Y = y|X = x) \\ &= - \sum_x P(x) \sum_y H(Y|X = x) \end{aligned}$$


For each possible value of X ,
1. compute its probability
2. compute conditional entropy

- $H(Y|X=x)$ is the *specific conditional entropy*, i.e. the entropy of Y knowing the value of a specific attribute x .

Back to the credit risk example

$$\begin{aligned}
 H(Y|X) &\equiv - \sum_x P(x) \sum_y P(y|x) \log_2 P(y|x) \\
 &= - \sum_x P(x) \sum_y P(Y = y|X = x) \log_2 P(Y = y|X = x) \\
 &= - \sum_x P(x) \sum_y H(Y|X = x)
 \end{aligned}$$

$$H(\text{defaulted} | < 2\text{years} = \text{N}) = -\frac{4}{4+2} \log_2 \frac{4}{4+2} - \frac{2}{6} \log_2 \frac{2}{6} = 0.9183$$

$$H(\text{defaulted} | < 2\text{years} = \text{Y}) = -\frac{3}{4} \log_2 \frac{3}{4} - \frac{1}{4} \log_2 \frac{1}{4} = 0.8133$$

$$H(\text{defaulted} | < 2\text{years}) = \frac{6}{10} 0.9183 + \frac{4}{10} 0.8133 = 0.8763$$

Average entropy given value of "<2years" attribute

$$H(\text{defaulted} | \text{missed} = \text{N}) = -\frac{6}{7} \log_2 \frac{6}{7} - \frac{1}{7} \log_2 \frac{1}{7} = 0.5917$$

$$H(\text{defaulted} | \text{missed} = \text{Y}) = -\frac{1}{3} \log_2 \frac{1}{3} - \frac{2}{3} \log_2 \frac{2}{3} = 0.9183$$

$$H(\text{defaulted} | \text{missed}) = \frac{7}{10} 0.5917 + \frac{3}{10} 0.9183 = 0.6897$$

Average entropy given value of "missed" attribute

Predicting credit risk

<2 yrs	missed	def?
N	N	N
Y	N	Y
N	N	N
N	N	N
N	Y	Y
Y	N	N
N	Y	N
N	Y	Y
Y	N	N
Y	N	N

- We now have the entropy - the minimal number of bits required to specify the target attribute:

$$H(Y) = \sum_y P(y) \log_2 P(y)$$

- The conditional entropy - the remaining entropy of Y knowing X

$$H(Y|X) = - \sum_x P(x) \sum_y P(y|x) \log_2 P(y|x)$$

- So we can now define the reduction of the entropy after learning Y.
- This is known as the *mutual information* between Y and X

$$I(Y; X) = H(Y) - H(Y|X)$$

Original
Entropy

Average entropy after
sorting on attribute X

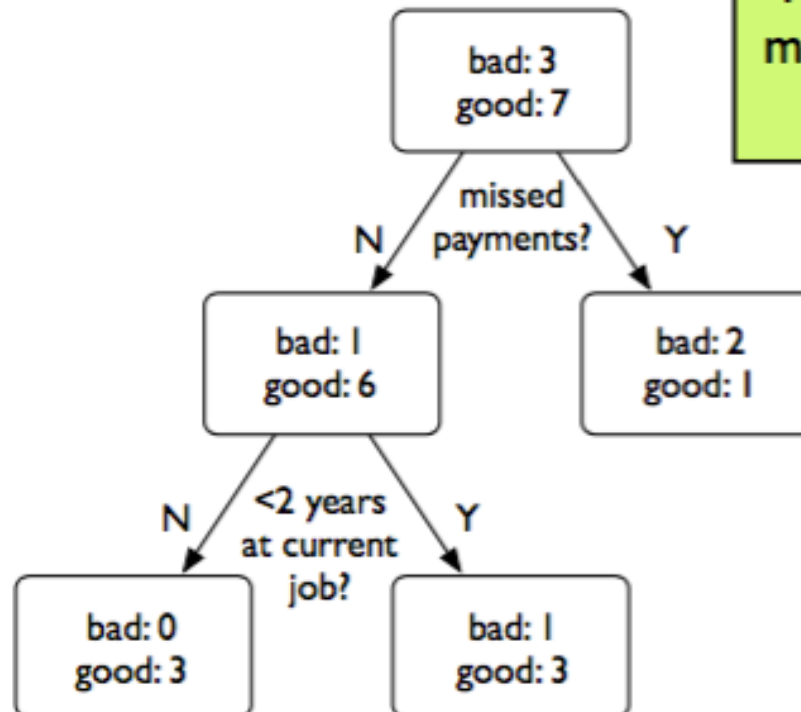
So... which attribute should I split on?



$$\begin{array}{rcl} H(\text{defaulted}) & - & H(\text{defaulted} | < 2 \text{ years}) \\ 0.8813 & - & 0.8763 = 0.0050 \end{array}$$

$$\begin{array}{rcl} H(\text{defaulted}) & - & H(\text{defaulted} | \text{missed}) \\ 0.8813 & - & 0.6897 = 0.1916 \end{array}$$

Missed payments are the most informative attribute about defaulting.




Calculating Information Gain


$$\text{Information Gain} = \text{entropy}(\text{parent}) - [\text{average entropy}(\text{children})]$$

$$\text{child entropy} = -\left(\frac{13}{17} \cdot \log_2 \frac{13}{17}\right) - \left(\frac{4}{17} \cdot \log_2 \frac{4}{17}\right) = 0.787$$

Entire population (30 instances)




Sorting on some attribute



17 instances

$$\text{child entropy} = -\left(\frac{1}{13} \cdot \log_2 \frac{1}{13}\right) - \left(\frac{12}{13} \cdot \log_2 \frac{12}{13}\right) = 0.391$$

$$\text{parent entropy} = -\left(\frac{14}{30} \cdot \log_2 \frac{14}{30}\right) - \left(\frac{16}{30} \cdot \log_2 \frac{16}{30}\right) = 0.996$$



13 instances

$$\text{(Weighted) Average Entropy of Children} = \left(\frac{17}{30} \cdot 0.787\right) + \left(\frac{13}{30} \cdot 0.391\right) = 0.615$$

$$\text{Information Gain} = 0.996 - 0.615 = 0.38 \quad \text{for this split}$$

Calculating IG

$$E(S) = E(29, 35)$$

$$E(X1) = E(21, 5)$$



$$E(X2) = E(30, 8)$$

$$IG = E(S) - [26/64 * E(X1) + 38/64 * E(X2)]$$

Information Gain

$Gain(S, A) =$ expected reduction in entropy due to sorting on A

$$Gain(S, A) \equiv Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$



ID3 algorithm

Top-Down Induction of Decision Trees

[ID3, C4.5, Quinlan]

node = Root

Main loop:

1. $A \leftarrow$ the “best” decision attribute for next *node*

one that gives
the best IG

2. Assign A as decision attribute for *node*

3. For each value of A , create new descendant of *node*

4. Sort training examples to leaf nodes

5. If training examples perfectly classified, Then
STOP, Else iterate over new leaf nodes


Worked out example

- Please see the handout/notes for a worked out example using ID3
- Remember:
 - For each node, you have to find the best attribute
 - You can only use an attribute once along a path. So, a node needs to inherit a list of attributes from its parent class
 - > You have to program this. 😊
 - At the leaf node, find the majority class (by count). Use that for the prediction rule.

Does it really matter which attribute comes first?



- ID3 helps us in selecting the shortest i.e. most compact tree



- ID3 performs heuristic search through space of decision trees
- It stops at smallest acceptable tree. Why?

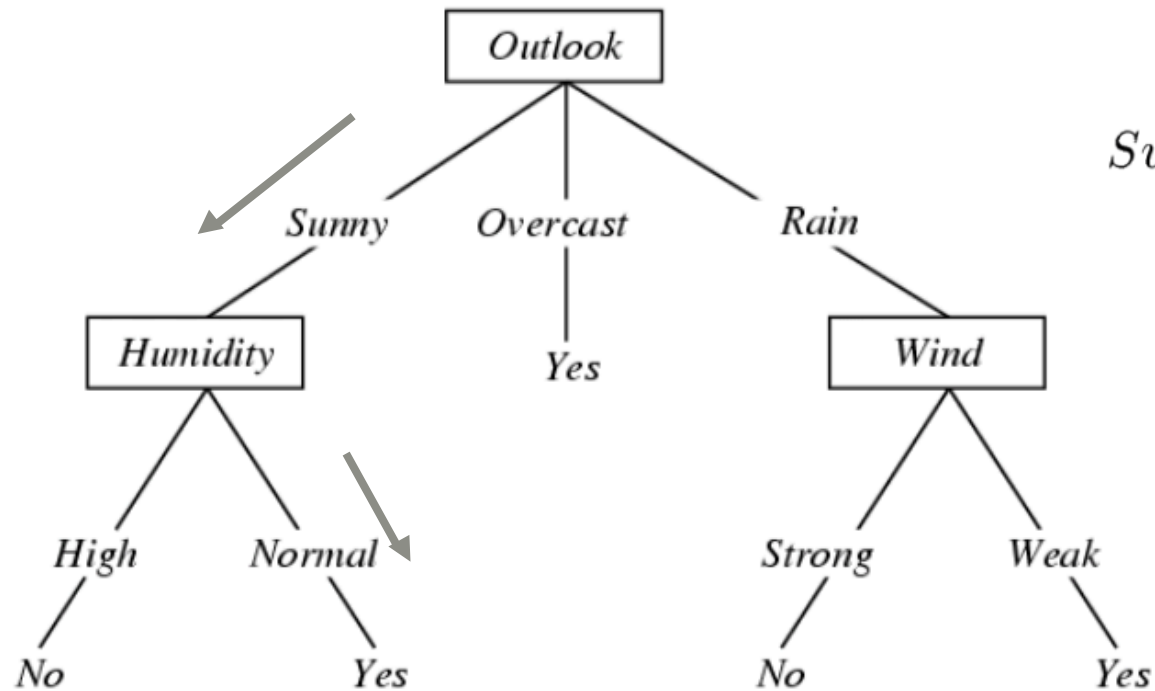
Occam's razor: prefer the simplest hypothesis that fits the data

ID3 is a greedy algorithm
Top-down induction of trees

Problem with DT

Problem with DT?

- Over-zealous learner -> learns all features
- What if there is noise?
- DT will try to change everything??
- Consider tree below:



What would happen if you get **noisy** data point

Sunny, Hot, Normal, Strong, PlayTennis = No

Overfitting

- You train on the training dataset

The data that the learner trains with.

=> It is possible to design a DT that gives 100% accuracy on training data. Think how??
e.g. each instance gets its own leaf node

- But is that a good thing?

=> NO! Because you are in fact memorizing (rote learning) the training data

=> No room for generalization, it's a case of Overfitting

- So, you have to find a balance between underfitting (learning very little) and overfitting.

- Notation:

Training error of hypothesis $h = e_{\text{train}}(h)$

True error (on unseen data) of hypothesis $h = e_{\text{true}}(h)$

Overfitting

Consider a hypothesis h and its

- Error rate over training data: $error_{train}(h)$
- True error rate over all data: $error_{true}(h)$


We say h overfits the training data if

$$error_{true}(h) > error_{train}(h)$$

Amount of overfitting =

$$error_{true}(h) - error_{train}(h)$$

Overfitting in Decision Tree Learning



How to avoid overfitting?

1. Stop growing when splits are not statistically significant (TOUGHER PROBLEM)

OR

2. Grow full tree then post-prune i.e. remove nodes and see if true error decreases (EASIER PROBLEM)

How to avoid overfitting?

- Keep another dataset -> validation dataset
- Build model on training, test accuracy on validation
- Learn model from training dataset.
- Randomly remove nodes and see if validation accuracy improves

Reduced-Error Pruning

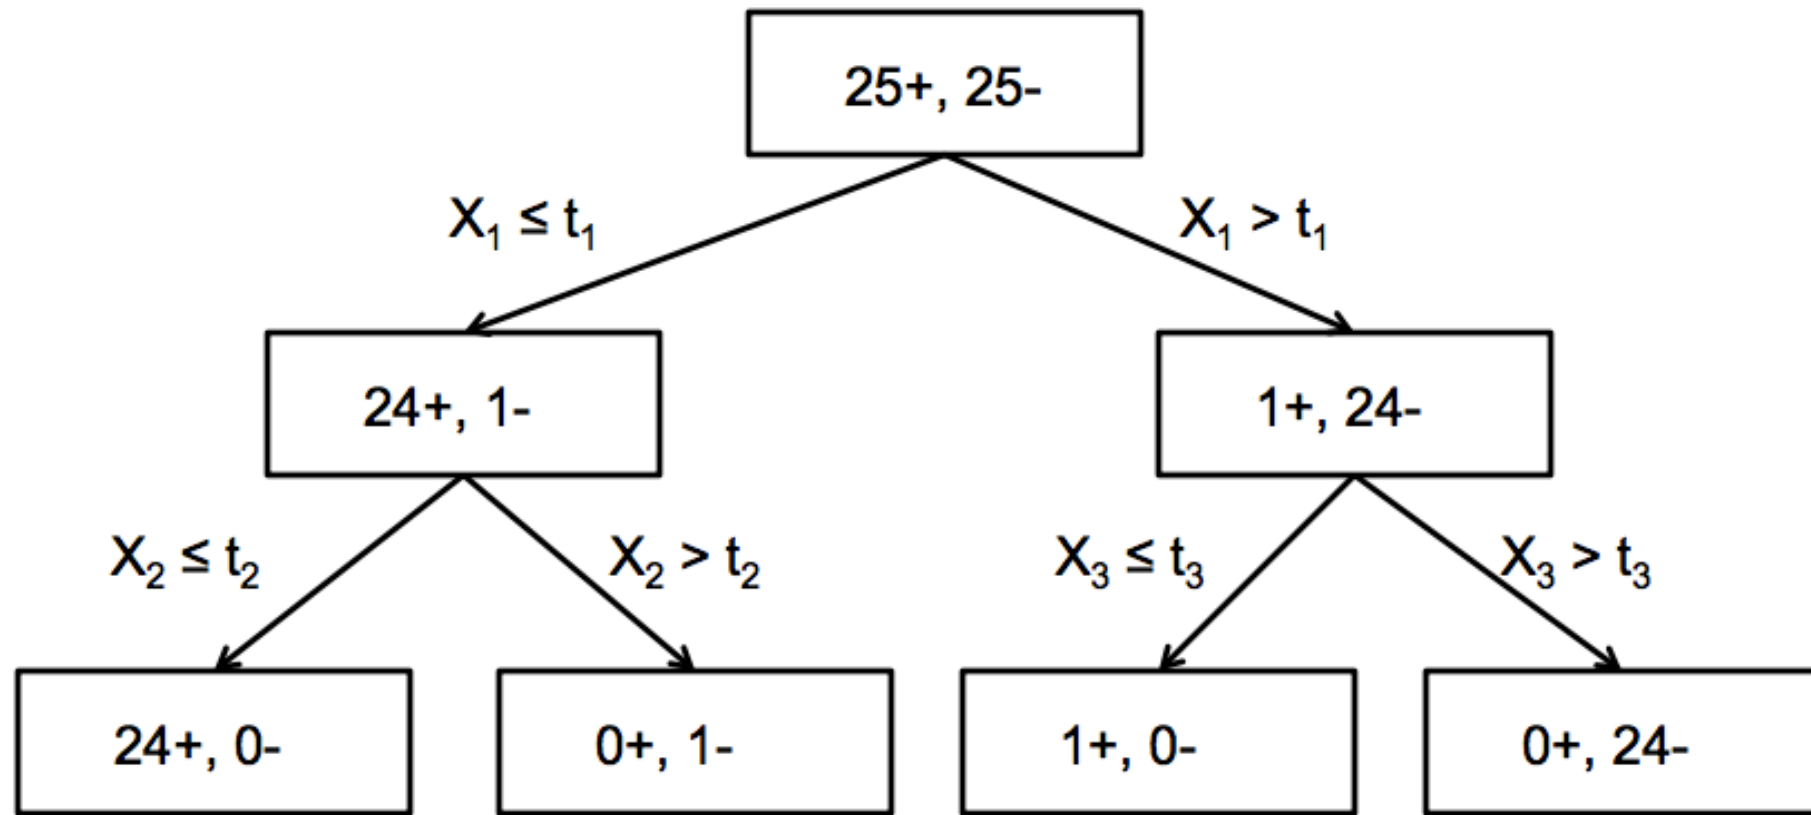
Split data into *training* and *validation* set

Create tree that classifies *training* set correctly

Do until further pruning is harmful:

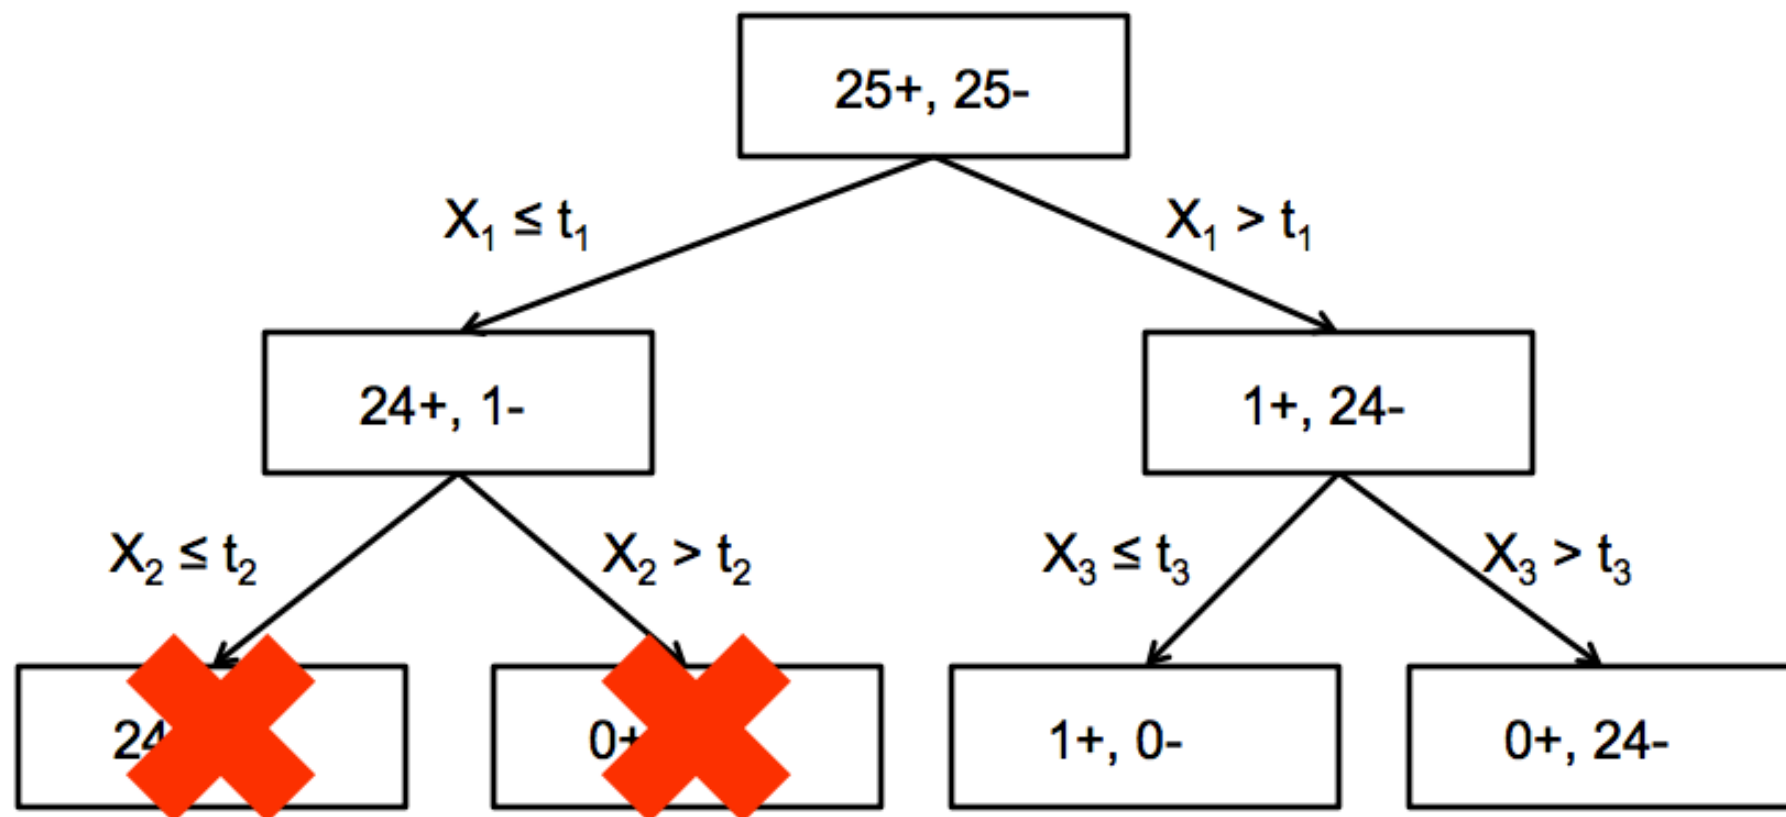
1. Evaluate impact on *validation* set of pruning each possible node (plus those below it)
2. Greedily remove the one that most improves *validation* set accuracy

Post-pruning



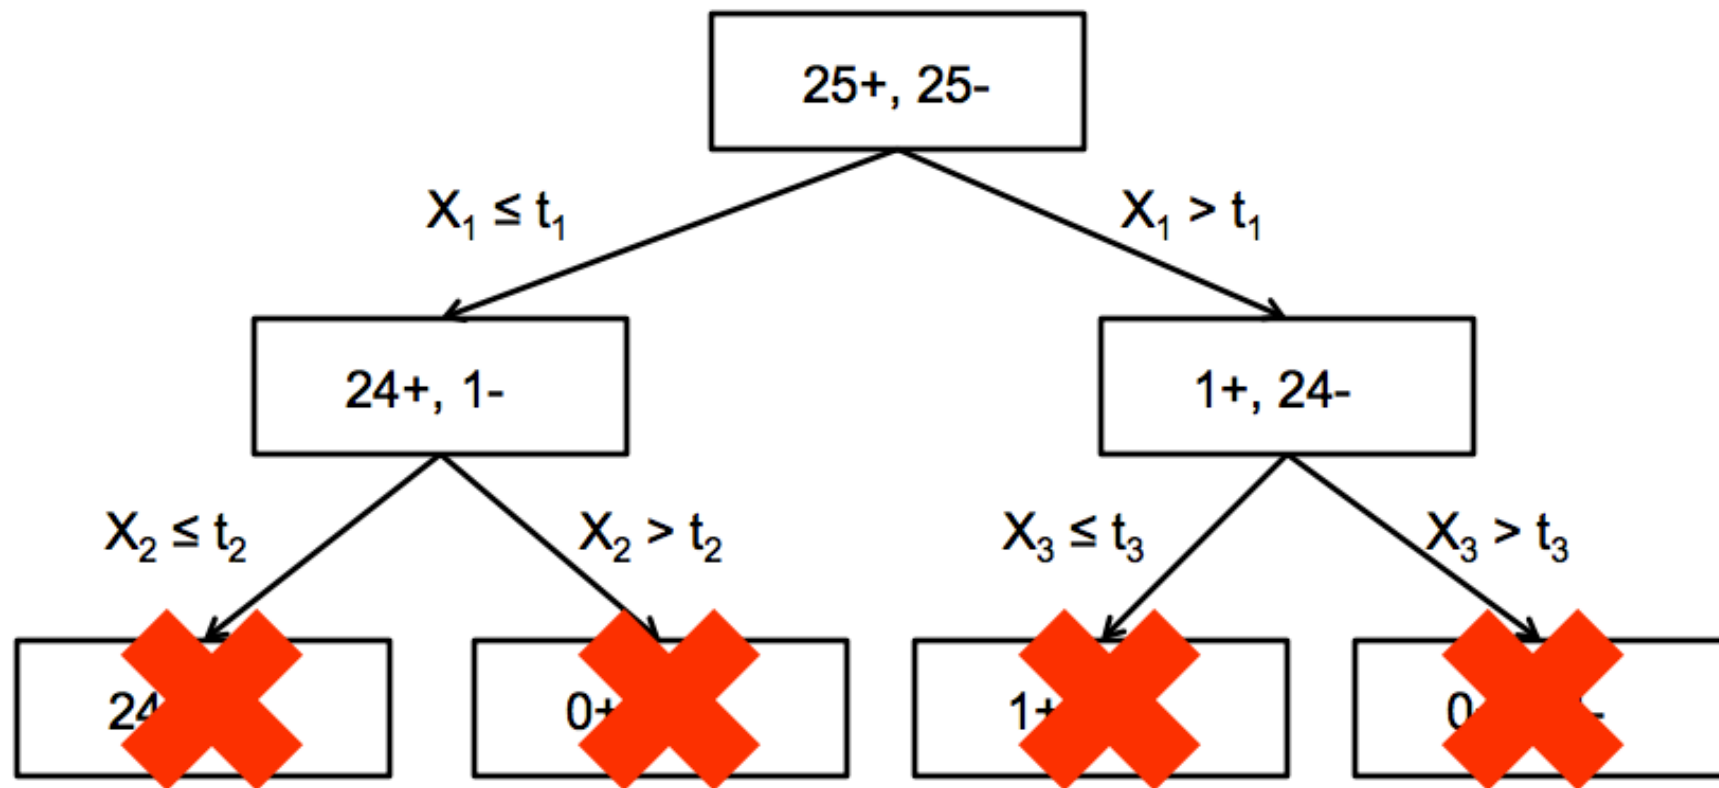
Validation Set Accuracy: 80%

Post-pruning



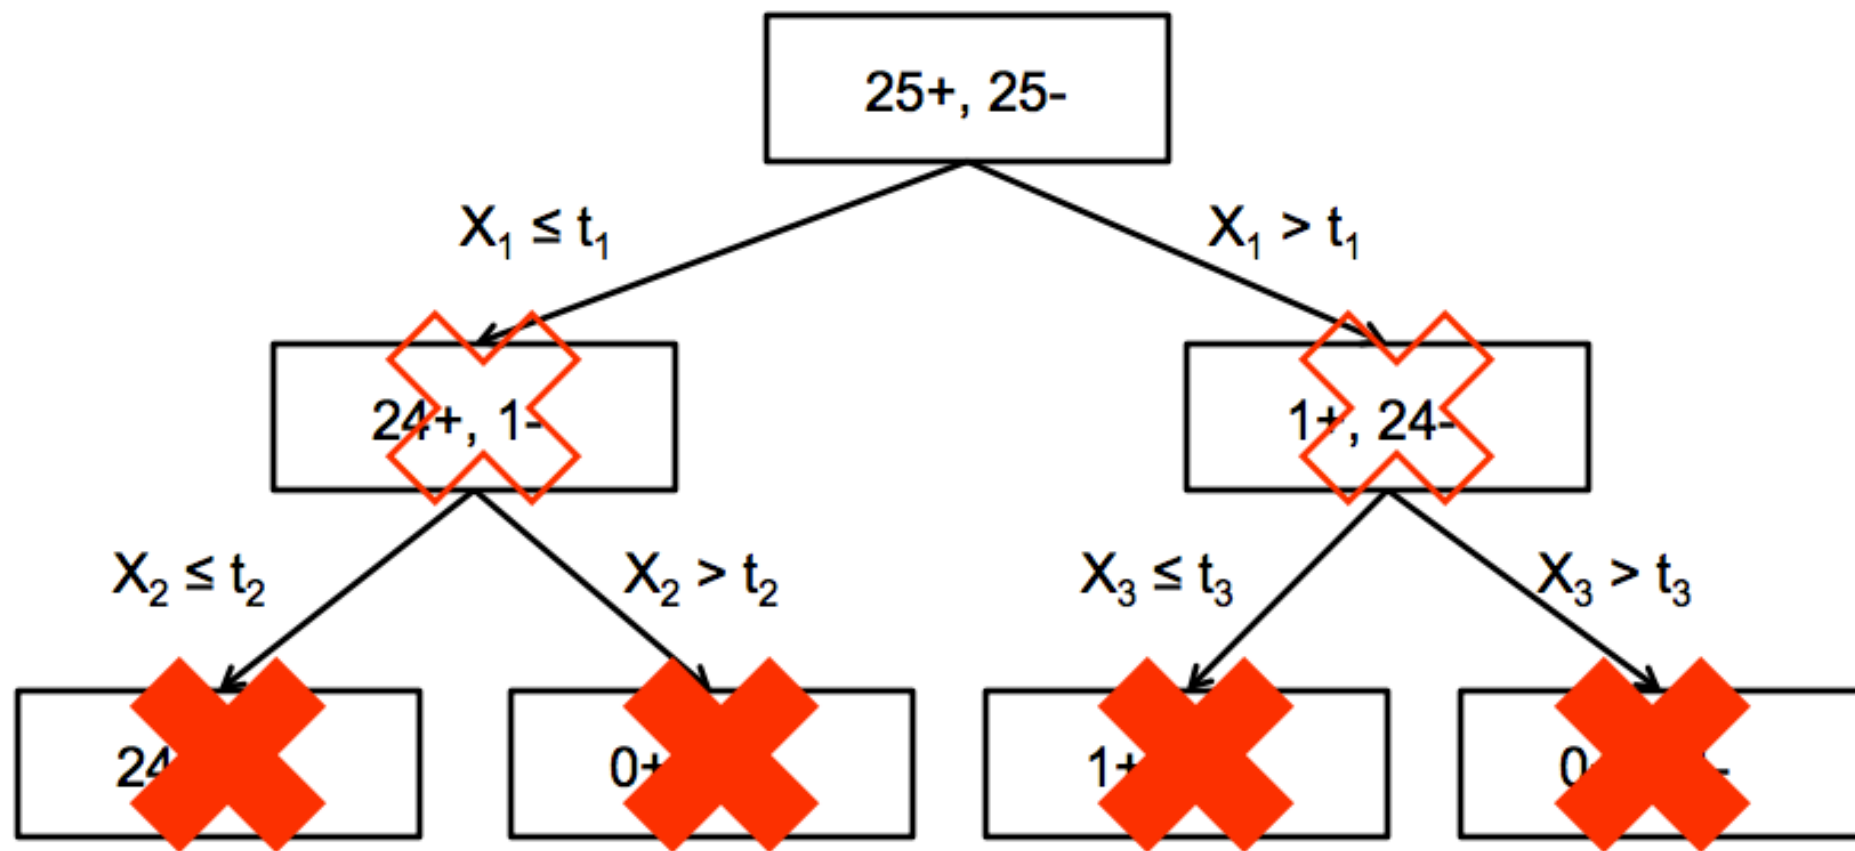
Validation Set Accuracy: 85%

Post-pruning




Validation Set Accuracy: 90%

Post-pruning



Validation Set Accuracy: 50%

Post-pruning



Final Decision Tree

What have we learnt?

- Idea of DT
- Number of instances (leaf nodes) and hypotheses
- How to choose best sorting attribute for each node
- How to induce top-down tree using ID3
- What is overfitting
- Avoiding overfitting