

Instance Based Learning

Review of Types of Classifiers

We have studied 3 types of classifiers:

1. Create a model for y (output) as a function of attributes. e.g. Perceptron, ANN, SVM

$$y = \text{sign}(w^T x)$$


2. Probabilistic (Generative) classifiers

$$P(Y | X) \propto P(X | Y) * P(Y)$$

Probability of class membership estimated from likelihood and prior in the training data.

3. Discriminative classifiers –

Create a model for $P(Y|X)$ – Logistic Regression

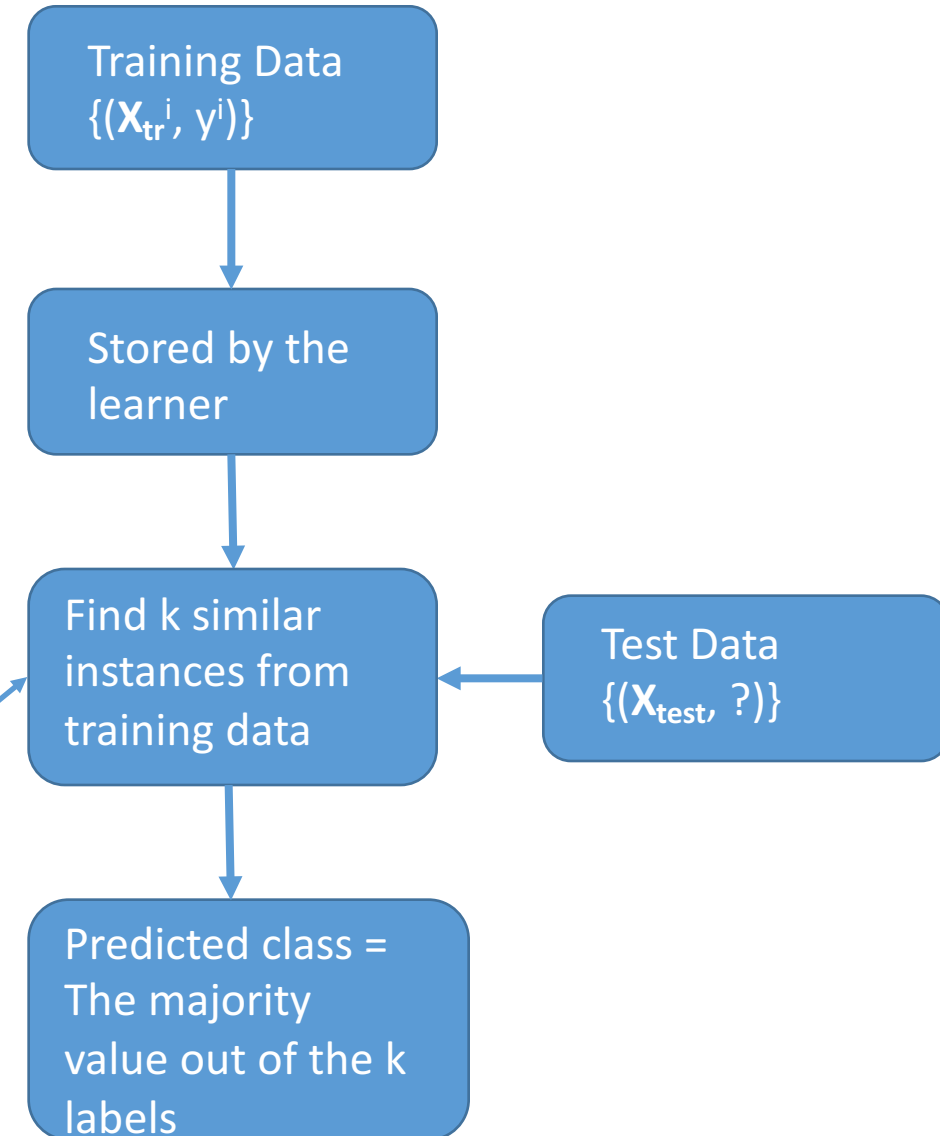


Today's topics

- Instance Based Learning
 - k-Nearest Neighbors
 - Locally Weighted Regression (LWR)
- Practice Questions
- Reading:
Chapter 8 of Tom Mitchell

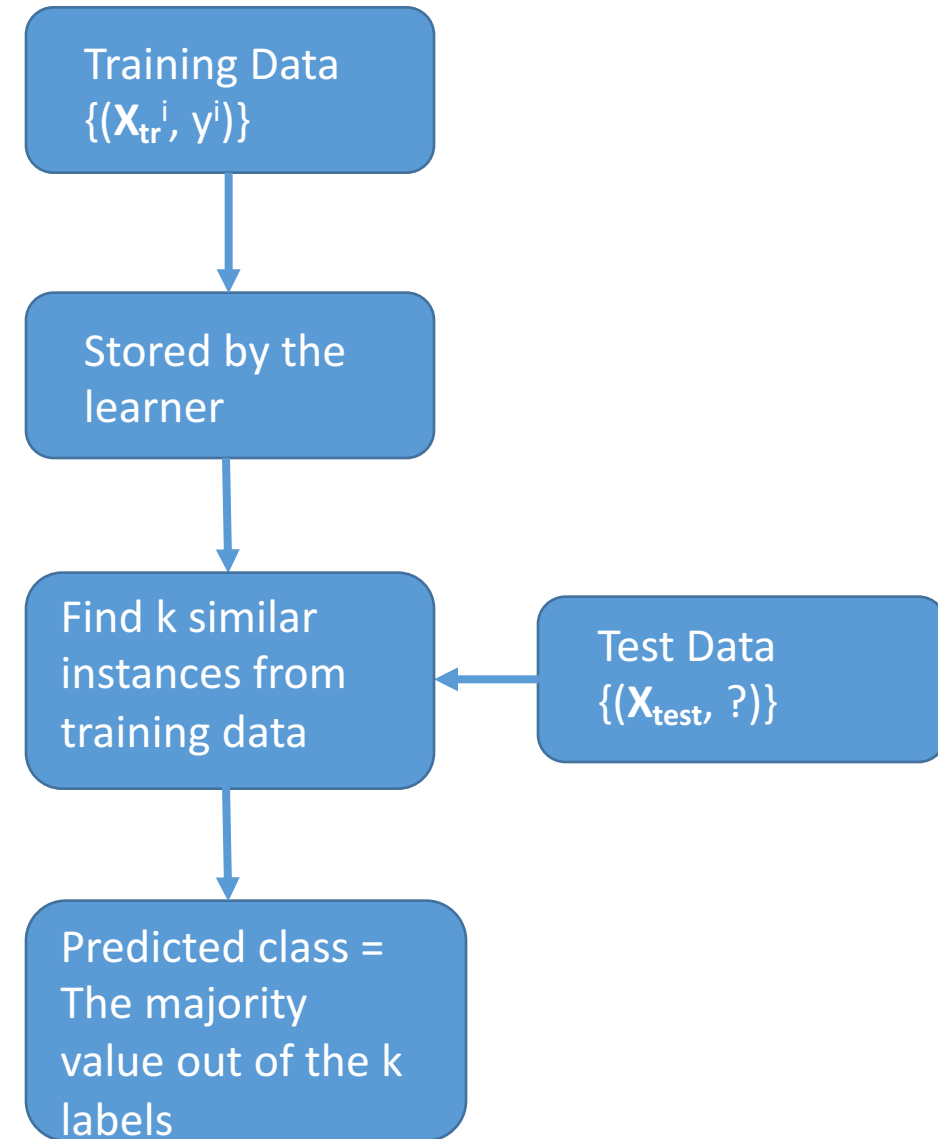
Instance Based Learning - Classification

- Learning phase consists of simply storing all the training examples [referred to as **instances**]
- In fact there is no model creation, just storing the data.
- When a query (test) data is presented, a set of similar instances is retrieved from memory and used to classify the data.
- k-Nearest Neighbors (kNN) is the most popular IBL algorithm.



k- Nearest Neighbors (kNN)

- What happens during testing phase?
 - **k** similar items are scanned
 - their class labels are identified
 - majority of the class labels gives the predicted class label.
- What are the parameters?
 - Similarity metric
 - Value of k



k-NN Algorithm

Training data: Labeled instances $\{\mathbf{X}^i, y^i\}$

Training Phase: Just store the examples

Test Phase:


Given a test instance \mathbf{X}^q , take a vote among its k-nearest neighbors for the class value:

$$\hat{f} = \text{majority}(f(X^i))$$

where $i \in \text{Neighbors}(K)$

Example


- Image on the right shows a k-NN example with $k = 1$ and $k = 5$.



Decision Boundary

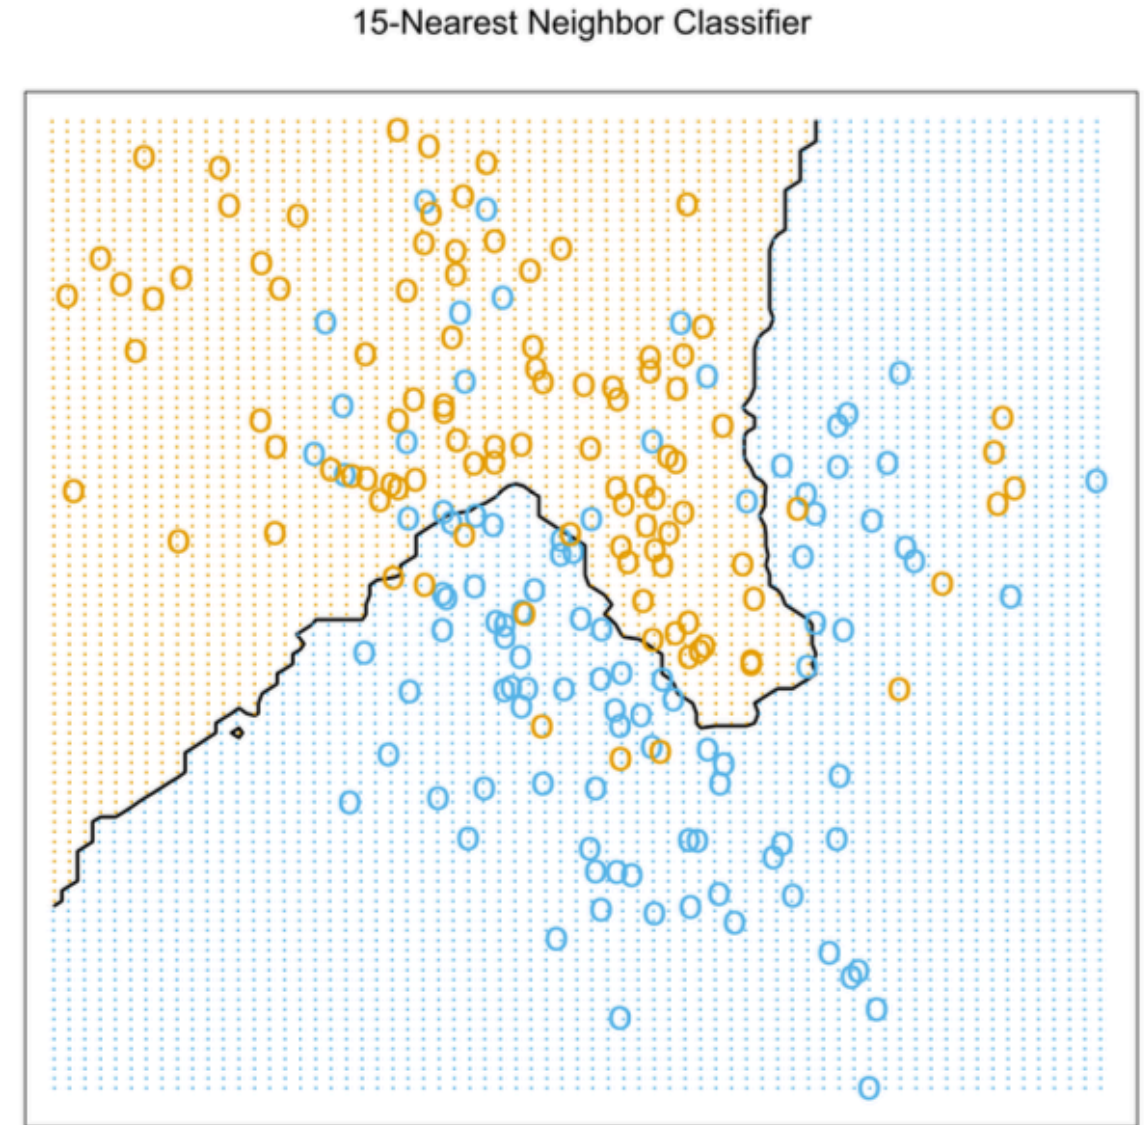
- The decision boundary for a 1-NN case can be shown as:
- This forms a Voronoi diagram,
Defn: partitioning of a plane into part based on distance to specific points.

1-NN Decision Surface




Example

- Image on the right shows a k-NN example with $k = 15$.




Example

- Image on the right shows a k-NN example with $k = 1$.
- Point to note:
You can get really complex decision boundaries by this very simple model.
- It pays to be lazy sometimes.



1-Nearest Neighbor Classifier



Distance Metric

- k-NN is critically dependent on the distance metric.
- For a function D to be used as a distance metric, it should satisfy certain properties:
 - ❑ $D(a, b) \geq 0$ (Positivity)
 - ❑ $D(a, b) = 0$ iff $a = b$ (Reflexive)
 - ❑ $D(a, b) = D(b, a)$ (Symmetric)
 - ❑ For any other vector c ,
 $D(a, b) + D(b, c) \geq D(a, c)$ (Triangle Inequality)
- What are some of the commonly used distance metrics?

Similarity Metrics

Some commonly used similarity metrics for vectors \mathbf{X}^i and \mathbf{X}^j having d dimensions each:

- Euclidean distance:

$$D(\mathbf{X}^i, \mathbf{X}^j) = \left[\sum_{a=1}^d (x_a^i - x_a^j)^2 \right]^{1/2}$$

- Manhattan distance:

$$D(\mathbf{X}^i, \mathbf{X}^j) = \left[\sum_{a=1}^d |x_a^i - x_a^j| \right]$$

Example:

If $X1 = (2, -1)^T$ and $X2 = (-2, 2)^T$

Euclidean Distance (D)=

$$\left[(2 - (-2))^2 + (-1 - 2)^2 \right]^{1/2} = 5$$

Manhattan Distance (D)=

$$[|2 - (-2)| + |-1 - 2|] = 7$$

Similarity Metrics

Both of these are generalization of the Minkowski distance (or norm)

- Minkowski distance of order p:

$$L_p = D(\mathbf{X}^i, \mathbf{X}^j) = \left[\sum_{a=1}^d |x_a^i - x_a^j|^p \right]^{1/p}$$

- p can range from 1 to ∞ .

Trivia: p cannot be less than 1, otherwise it would violate triangle inequality.

It's fun exercise to show this. Try it on your own.

Hint: Set a = (0, 0) b = (1, 1) and c = (0, 1) and p = ½

Does it still satisfy the triangle inequality?

A word of caution

- For these distance metrics to work nicely, the attributes must be scaled before using them.
- We have done this earlier in neural net training.
- In many cases, you might want to weight each attribute differently. This is called weighted distance.
- Example of weighted Euclidean distance:

$$D_w(\mathbf{X}^i, \mathbf{X}^j) = \left[\sum_{a=1}^d w_a (x_a^i - x_a^j)^2 \right]^{1/2}$$

where w_a is weight for the a^{th} attribute

Example:

Data for UTD students' GPA (attribute) and getting internship (class) is presented below:

GPA	2.6	2.8	2.85	3.1	3.2	3.3	3.4	3.55	3.6	3.7	3.75	4.0	4.0
Internship	0	1	1	0	1	1	0	1	0	1	0	0	1

Using Manhattan distance, what will be the prediction of k-NN for a student with GPA of 3.5 in the following cases:

1. $k = 1$
2. $k = 3$
3. $k = 5$

Example:

Data for UTD students' GPA (attribute) and getting internship (class) is presented below:

GPA	2.6	2.8	2.85	3.1	3.2	3.3	3.4	3.55	3.6	3.7	3.75	4.0	4.0
Internship	0	1	1	0	1	1	0	1	0	1	0	0	1

Using Manhattan distance, what will be the prediction of k-NN for a student with GPA of 3.5 in the following cases:

1. **k = 1:** Nearest neighbor: {(3.55, 1)} => Majority class = 1
2. **k = 3:** Nearest neighbors: {(3.55, 1), (3.4, 0), (3.6, 0)} => Majority class = 0
3. **k = 5:** Nearest neighbors: {(3.55, 1), (3.4, 0), (3.6, 0), (3.3, 1), (3.6, 1)}
=> Majority class = 1

Handout

- Let's practice some questions from the handout.

Handout

- Let's practice some questions from the handout.
- Realization:
 - The calculations become more and more tedious as k increases and more so when dimensions increase.
 - More work is done during the testing phase than during the training phase.

Advantages and Disadvantages of k-NN

Advantages:

- Training is very fast
- Can learn complex target functions easily
- Easy to program

Disadvantages:

- Slow at query time
- Lots of storage and processing in memory
- Doesn't scale well in higher dimensions **Curse of Dimensionality**
- Easily tricked by noisy data items and irrelevant attributes
- Value of k can vary results significantly

k-NN for Continuous Output

- We presented k-NN for classification, but it can easily be used to approximate functions where the output is continuous.
- How? Simply replace the majority function by the average function:
Instead of:

$$\hat{f} = \text{majority}(f(X^i))$$

where $i \in \text{Neighbors}(K)$


Use this:

$$\hat{f} = \frac{\sum_{i=1}^k f(x_i)}{k}$$


where $i \in \text{Neighbors}(K)$

Locally Weighted Local Regression

- In typical regression, we create a global model for the data.
- A better option? Creating local approximations using the idea of "neighborhood" and k-NN.
- The local functions produce piecewise approximation to the global function.



LWR Example



Training data



Predicted value using simple regression



Predicted value using locally weighted (piece-wise) regression

Lazy vs Eager Learners

- **Lazy:** Wait for query before generalizing i.e. going beyond training data.
 - > Model doesn't propose a function until a query point arrives.
 - > When a query point arrives, a local search for the best function is done and it is outputted as the generalization for the test query.
 - > Model is adaptable.
- **Eager:** After all training data is received, a generalization function is outputted.
 - > ID3 algorithm creates a decision tree (generalization model) after incorporating all training data.
 - > A generalized linear regression outputs an approximation function after seeing all the training data.
 - > One model for all the query instances.

Parametric vs Non-Parametric Methods

- **Parametric:** A particular **functional form** is assumed, e.g., linear, Gaussian distribution, naïve Bayes
 - > may have **high bias** because the real data may not obey the assumed functional form.
 - > Parameter estimation is required.
 - > Advantage of simplicity – easy to estimate and interpret
- **Non-Parametric:** Distribution or density estimate is **data-driven** and relatively **few assumptions are made a priori about the functional form**.
 - > Parameter estimation is not as involved.

Thought Questions

- Can you think of a lazy algorithm for decision tree construction?
- Will it help if we are able to pre-process and remove irrelevant or redundant attributes?
- What is the time complexity of k-NN?